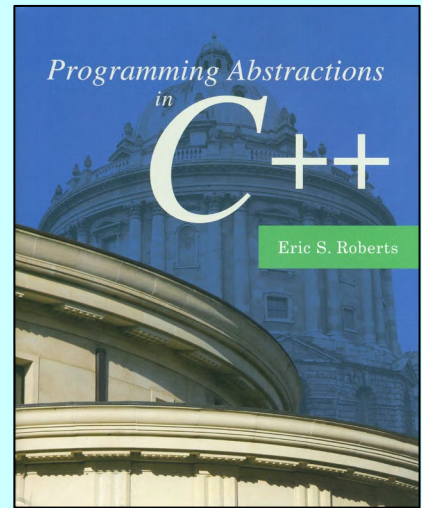


## CHAPTER 16

# Trees

*I like trees because they seem more resigned to the way they have to live than other things do.*

—Willa Cather, *O Pioneers!*, 1913



16.1 Family trees

16.2 Binary search trees

16.3 Balanced trees

16.4 Implementing maps using BSTs

16.5 Partially ordered trees

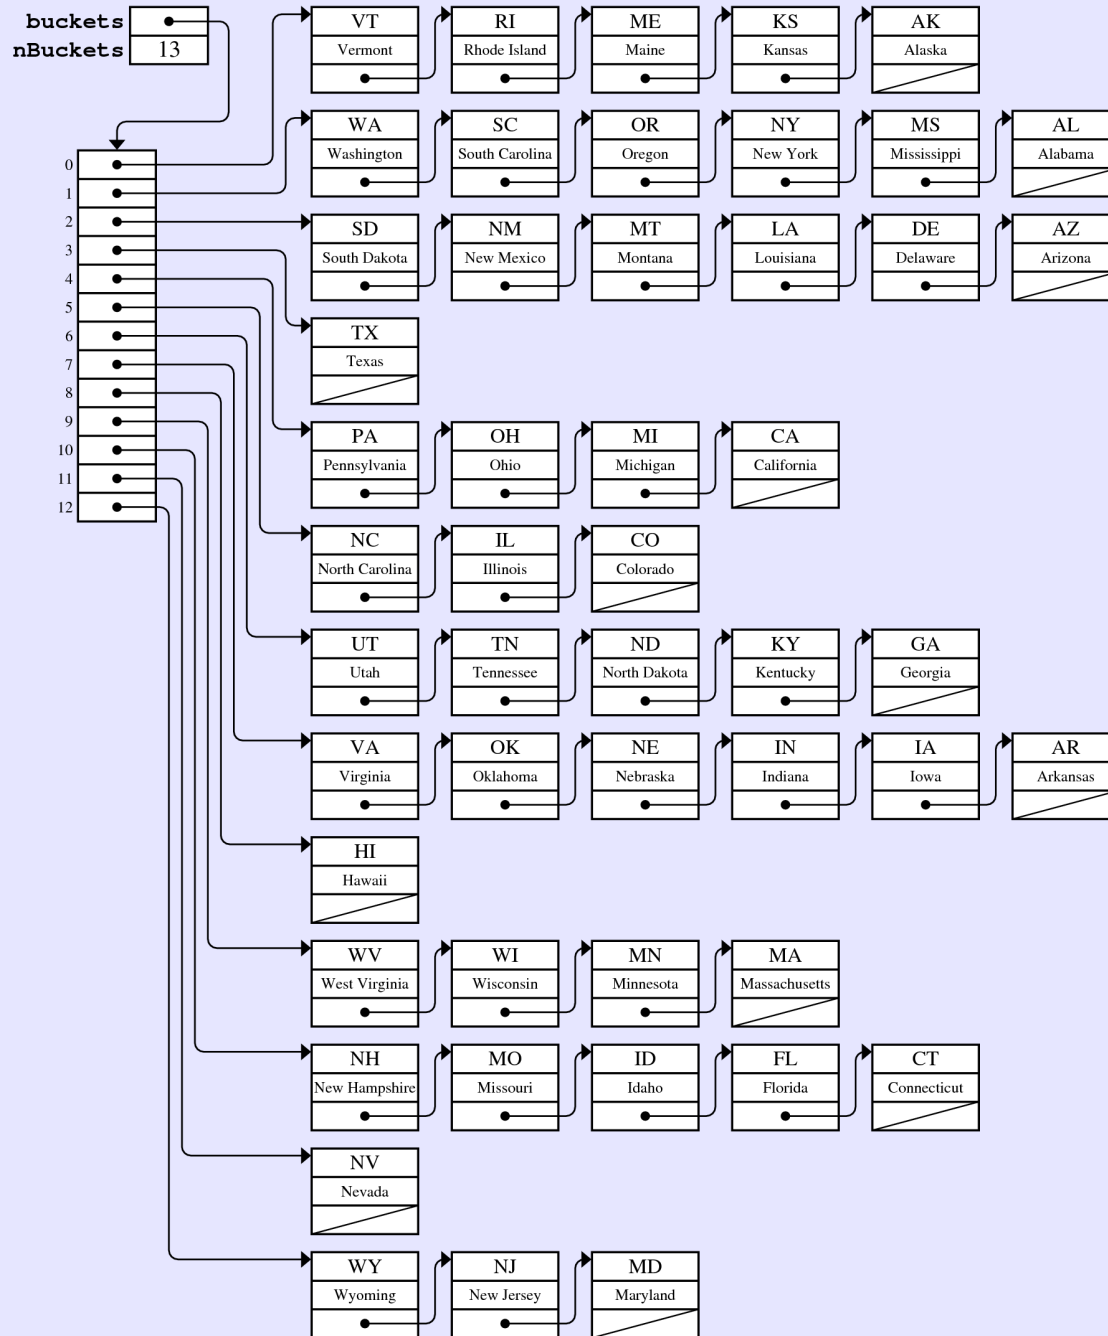
# Implementation Strategies for Maps

- There are several strategies you might choose to implement the map operations **get** and **put**. Those strategies include:
  1. **Linear search**. Keep track of all the key/value pairs in a vector. In this model, both the **get** and **put** operations run in  $O(N)$  time.
  2. **Binary search**. If you keep the vector sorted by the key, you can use binary search to find the key. Using this strategy improves the performance of **get** to  $O(\log N)$  and **put** is still  $O(N)$ .
  3. **Table lookup in a grid**. In the two-character code example, you can store the state names in a  $26 \times 26$  **Grid<string>** in which the first and second indices correspond to the two letters in the code. Because you can now find any code in a single step, this strategy is  $O(1)$ , although this performance comes **at a cost in memory space** (only 50/676 occupied).
  4. **Hashing**. Use a preferably  $O(1)$  time hash function to transform keys into as uniformly as possible distributed integer hash codes, which tell the implementation where it should look for a particular key, thereby reducing the search time dramatically.

# Limitations of Hashing

- In the last chapter, We have seen how hashing makes it possible to implement the **get** and **put** operations for a map in  $O(1)$  time.
- Despite its extraordinary efficiency, hashing is not always the best strategy for implementing maps, because of the following limitations:
  - Hash tables depend on being able to compute a **hash function** on some key. Expanding the hash-function idea so that it applies to types other than strings is subtle.
  - Using the range-based **for** on hash tables does not deliver the keys in any **sensible order**. Even when the keys have a natural order (such as the lexicographic order used with strings), the hash table code for **iteration** cannot take advantage of that fact.
- The goal now is to explore another representation that supports **iterating through the elements in order**.

**FIGURE 15-9** Hash table containing the state abbreviations





# Iterator Order

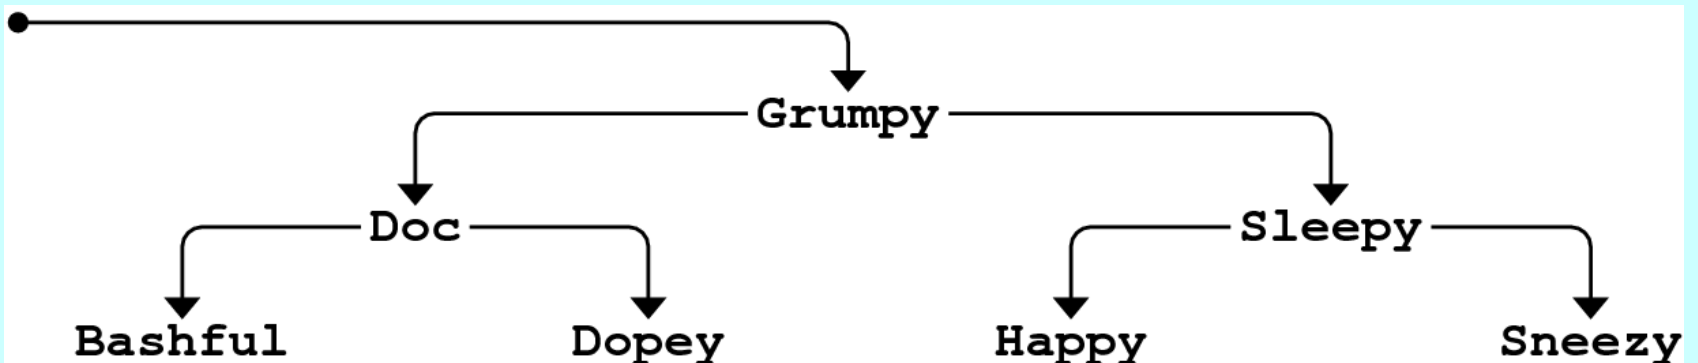


- When you look at the documentation for an iterator, one of the important things to determine is whether the collection class specifies the order in which elements are generated. The Stanford libraries make the following guarantees:
  - Iterators for the **vector** class operate in index order.
  - Iterators for the **Grid** class operate in *row-major order*, which means that the iterator runs through every element in row 0, then every element in row 1, and so on.
  - Iterators for the **Map** class deliver the keys in the order imposed by the standard comparison function for the key type; iterators for the **HashMap** class return keys in a **seemingly random** order.
  - Iterators for the **set** class deliver the elements in the order imposed by the standard comparison function for the value type; the **HashSet** class is unordered.
  - Iterators for the **Lexicon** class always deliver words in alphabetical order.

# Analyzing the Failure of Sorted Arrays

- One of the strategies outlined previously for implementing a map was to use a **sorted array** to hold the key-value pairs. Given that representation, binary search made it possible to **find a key in  $O(\log N)$  time**.
- The problem with the sorted array strategy was that **inserting a new key required  $O(N)$  time to *maintain the order***.
- In the editor buffer, linked lists solved the insertion problem. Unfortunately, turning a sorted array into **a linked list makes it impossible to apply binary search** because there is no way to go to the middle element in  $O(1)$  time.
- But what if you could **point to the middle element in a linked list**? That question gives rise to a new structure called a ***tree***, which provides the key to implementing a map with  $O(\log N)$  performance for both the **get** and **put** operations.

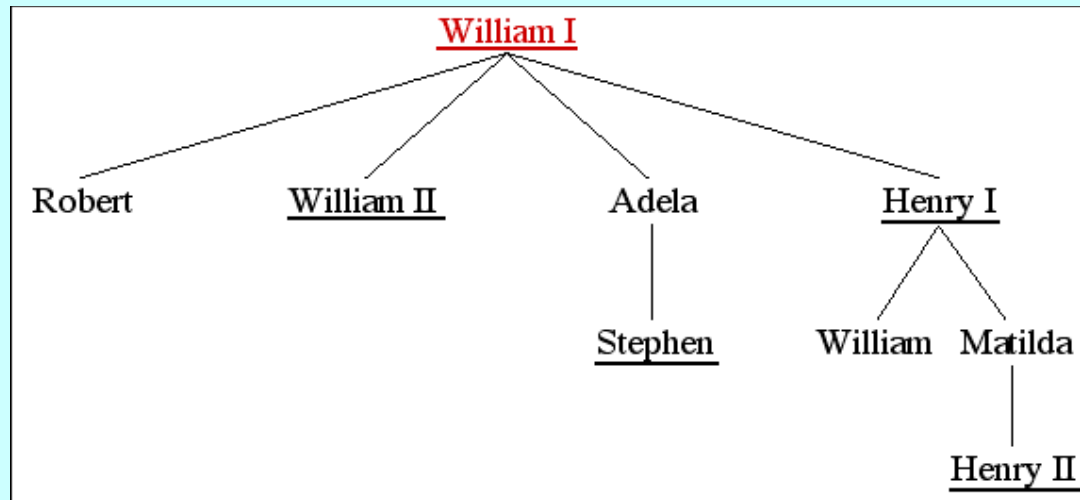
# The motivation behind Binary Search Trees

- The idea is to combine the **linked list** data structure (for easier insertion/deletion) and the **binary search** algorithm (for faster search).



# Trees

- In the text, the first example used to illustrate tree structures is the royal family tree of the House of Normandy:




- This example is useful for defining terminology:
  - William I is the *root* of the tree.
  - Adela is a *child* of William I and the *parent* of Stephen.
  - Robert, William II, Adela, and Henry I are *siblings*.
  - Henry II is a *descendant* of William I, Henry I, and Matilda.
  - William I is an *ancestor* of everyone else in this tree.



# Trees Are Everywhere

- Trees appear in many familiar contexts beyond family trees. The picture at the right comes from Darwin's notebooks and shows his early conception of an evolutionary tree.
- Trees also form the basis for the class hierarchies used in object-oriented programming languages like Python and C++.
- In each of these contexts, trees begin with a single root node and then branch outward repeatedly to encompass any other nodes in the tree.



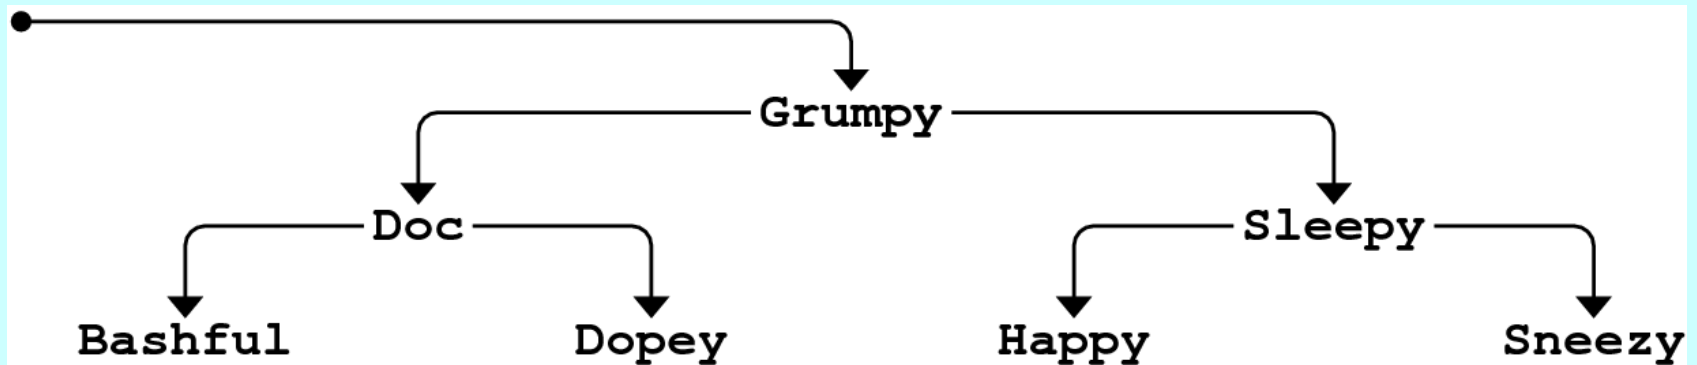
# Trees as a Recursive Data Structure

- If you think about trees as a programmer, the following definition is extremely useful:
  - A *tree* is a pointer to a node.
  - A *node* is a structure that contains some number of trees.
- Although this definition is clearly circular, it is not necessarily infinite, either because
  - Tree pointers can be **NULL** indicating an empty tree.
  - Nodes can contain an empty list of children.
- In C++, programmers typically define a structure or object type to represent a node and then use an explicit pointer type to represent the tree.

```
struct TreeNode {  
    string key;  
    Vector<TreeNode *> children;  
};
```

# Binary Search Trees

- The tree that supports the implementation of the **Map** class is called a *binary search tree* (or *BST* for short). Each node in a BST has exactly two subtrees: a *left subtree* that contains all the nodes that come before the current node and a *right subtree* that contains all the nodes that come after it. Either or both of these subtrees may be **NULL**.
- The classic example of a binary search tree uses the names from Walt Disney's *Snow White and the Seven Dwarves*:



# A Simple BST Implementation

- To get a sense of how binary search trees work, it is useful to start with a simple design in which keys are always strings.
- Each node in the tree is then a structure containing a key and two subtrees, each of which is either **NULL** or a pointer to some other node. This design suggests the following definition:

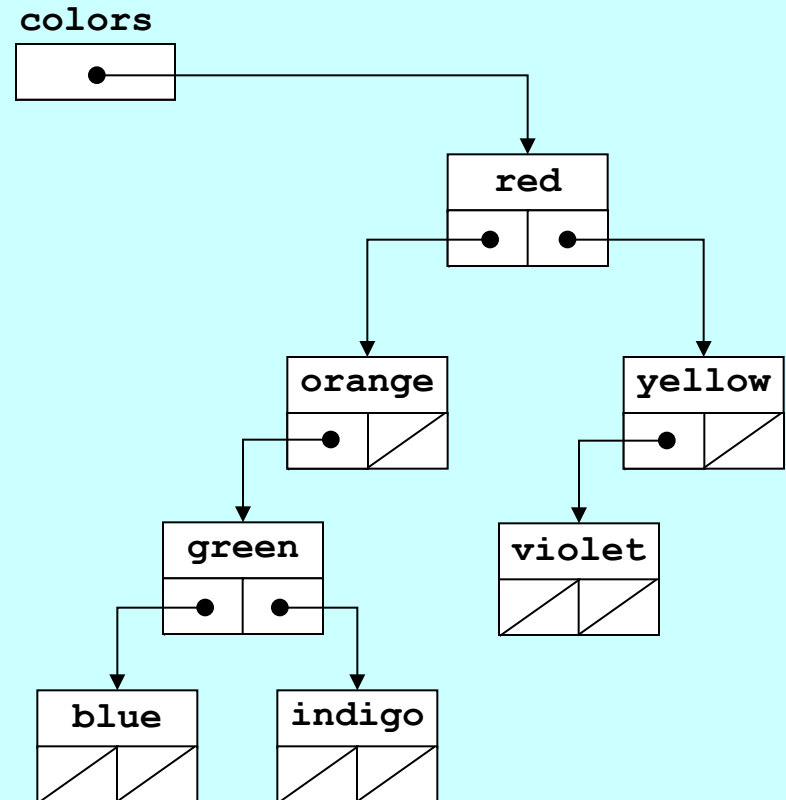
```
struct BSTNode {  
    string key;  
    BSTNode *left, *right;  
};
```

- The code for finding a node in a tree begins by comparing the desired key with the key in the root node. If the strings match, you've found the correct node; if not, you simply call yourself recursively on the left or right subtree depending on whether the key you want comes before or after the current one.

# Exercise: Building a Binary Search Tree

Diagram the BST that results from executing the following code:

```
Node *colors = NULL;  
insertNode(colors, "red");  
insertNode(colors, "orange");  
insertNode(colors, "yellow");  
insertNode(colors, "green");  
insertNode(colors, "blue");  
insertNode(colors, "indigo");  
insertNode(colors, "violet");
```



# A Simple BST Implementation

```
/*
 * Type: Node
 * -----
 * This type represents a node in the binary search tree.
 */

struct Node {
    string key;
    Node *left, *right;
};

/*
 * Function: findNode
 * Usage: Node* node = findNode(t, key);
 * -----
 * Returns a pointer to the node in the binary search tree t than contains
 * a matching key.  If no such node exists, findNode returns NULL.
 */

Node* findNode(Node* t, string key) {
    if (t == NULL) return NULL;
    if (key == t->key) return t;
    if (key < t->key) {
        return findNode(t->left, key);
    } else {
        return findNode(t->right, key);
    }
}
```

# A Simple BST Implementation

```
/*  
 * Function: insertNode  
 * Usage: insertNode(t, key);  
 * -----  
 * Inserts the specified key at the appropriate location in the  
 * binary search tree rooted at t. Note that t must be passed  
 * by reference, since it is possible to change the root.  
 */
```

```
void insertNode(Node* & t, string key) {
```

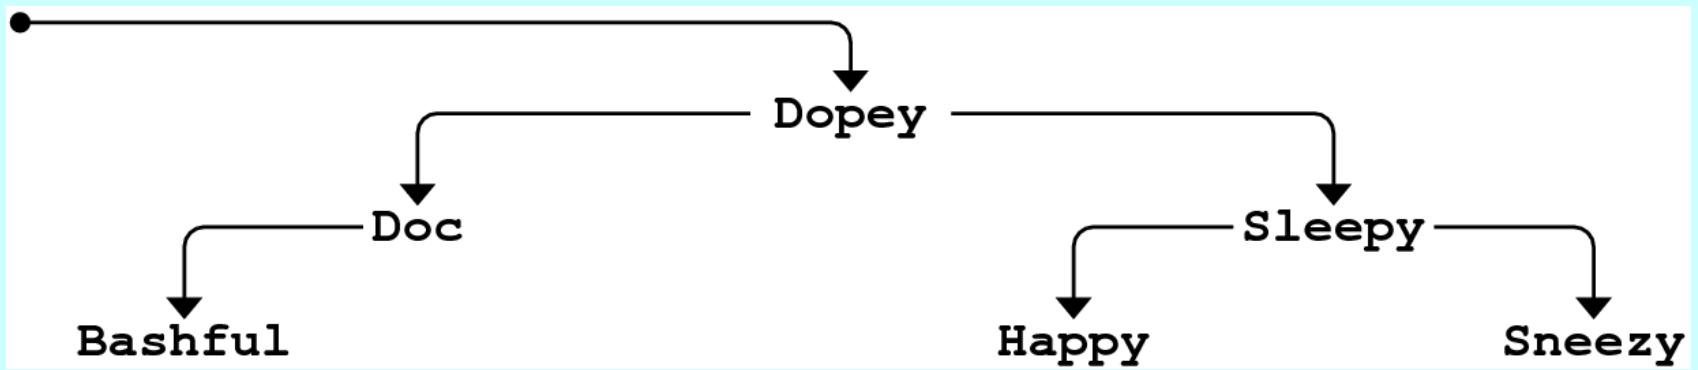
```
    if (t == NULL) {  
        t = new Node;  
        t->key = key;  
        t->left = t->right = NULL;  
        return;  
    }
```

*Call a pointer by reference. Can you use call by pointer here?*

```
    if (key == t->key) return;  
    if (key < t->key) {  
        insertNode(t->left, key);  
    } else {  
        insertNode(t->right, key);  
    }  
}
```

# Removing Nodes in Binary Search Trees

- For a leaf node (with no children), replace the pointer to the node with a NULL pointer;
- If either child of the node you want to remove is NULL, replace it with its non-NULL child;
- If you try to remove a node with both a left and a right child, replace it with the rightmost node in the left subtree (predecessor) or the leftmost node in the right subtree (successor). Additional operations might be needed.

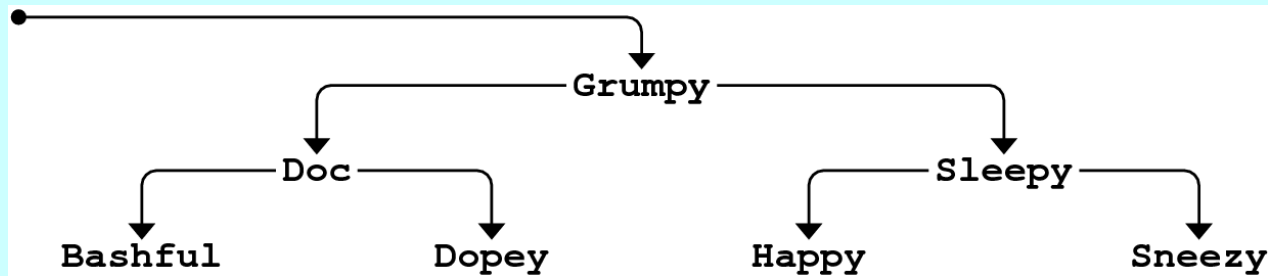






# A Question of Balance

- Ideally, a binary search tree containing the names of Disney's seven dwarves would look like this:






- If, however, you happened to enter the names in alphabetical order, this tree would end up being a simple linked list in which all the left subtrees were **NULL** and the right links formed a simple chain. Algorithms on that tree would run in  $O(N)$  time instead of  $O(\log N)$  time.






- A binary search tree is **balanced** if the height of its left and right subtrees differ by at most one and if both of those subtrees are themselves balanced.

**FIGURE 16-5** Examples of balanced and unbalanced binary trees

*Balanced trees:*




*Unbalanced trees:*




**FIGURE 16-7** Rotation operations in an AVL tree

*Single rotation*



*Double rotation*



*Things that you only need to know that you don't know. I don't need you to memorize it for the exam.*

Note: At least one of the subtrees  $T_2$  and  $T_3$  must have height  $h$ , the other can have height  $h$  or  $h-1$ . The balance factors in the final nodes will need to be adjusted to take account of any difference in height.

# Tree-Balancing Algorithms

- The AVL algorithm was the first tree-balancing strategy and has been superseded by newer algorithms that are more effective in practice. These algorithms include:
  - Red-black trees
  - 2-3 trees
  - AA trees
  - Fibonacci trees
  - Splay trees
- At this point in your study of computer science, the important thing to know is that it is *possible* to keep a binary tree balanced as you insert nodes, thereby ensuring that lookup operations run in  $O(\log N)$  time. If you get really excited about this kind of algorithms, you'll have the opportunity to study them in more detail in the later data structures and algorithms course.

# Implementing Maps Using BSTs

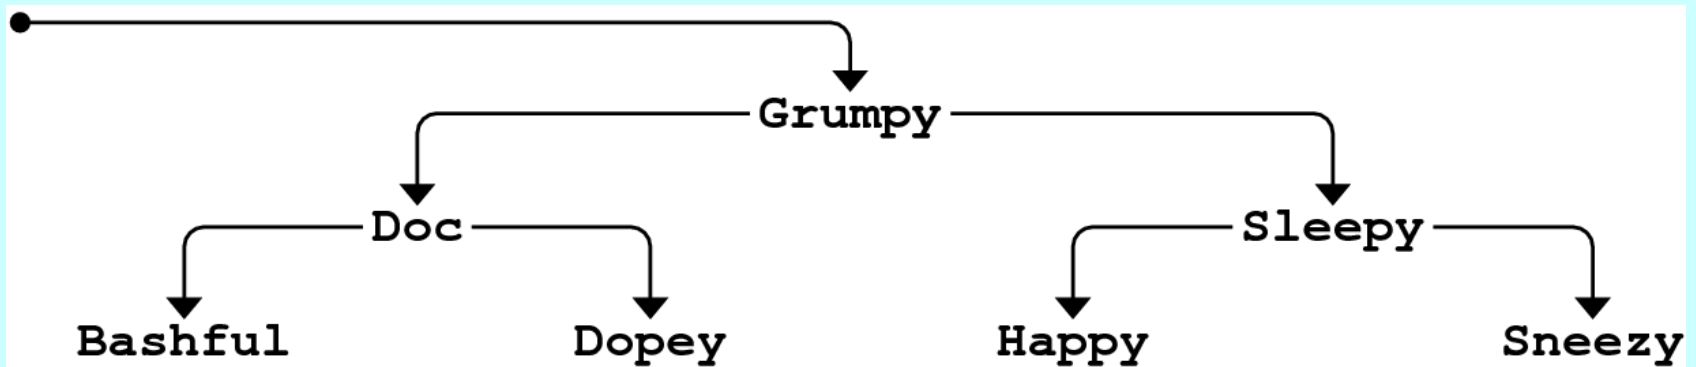
- The Stanford **Map** class (and **map** in Standard C++) uses a BST structure internally, in which **get** and **put** operate in  $O(\log N)$  time while allowing iteration to proceed in order.
- The following changes are required to implement the **Map** class:
  - The **BSTNode** structure must be expanded to include a **value** field.
  - The **Map** class must include template parameters for both the key and value types.
  - The implementation must include the methods that define the complete **Map** interface, such as **size**, **isEmpty**, **clear**, and **containsKey**.
  - The code for manipulating the tree must be embedded in the **Map** class.



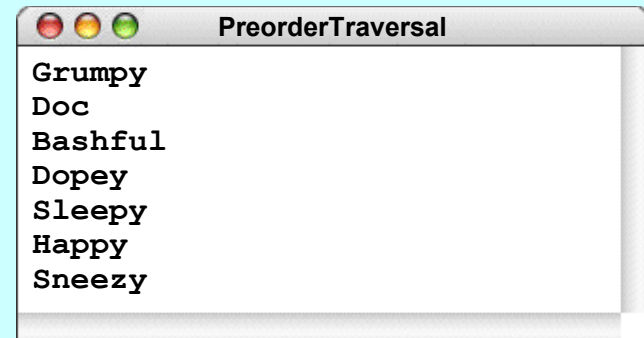
# Tree Traversal Strategies

- It is easy to write a function that performs some operation for every key in a binary search tree, because **recursion** makes it simple to apply that operation to each of the subtrees.
- The order in which keys are processed depends on when you process the current node with respect to the recursive calls:
  - If you process the current node before either recursive call, the result is a ***preorder traversal***.
  - If you process the current node after the recursive call on the left subtree but before the recursive call on the right subtree, the result is an ***inorder traversal***. In the case of the simple BST implementation that uses strings as keys, the keys will appear in lexicographic order.
  - If you process the current node after completing both recursive calls, the result is a ***postorder traversal***. Postorder traversals are particularly useful if you are trying to free all the nodes in a tree.

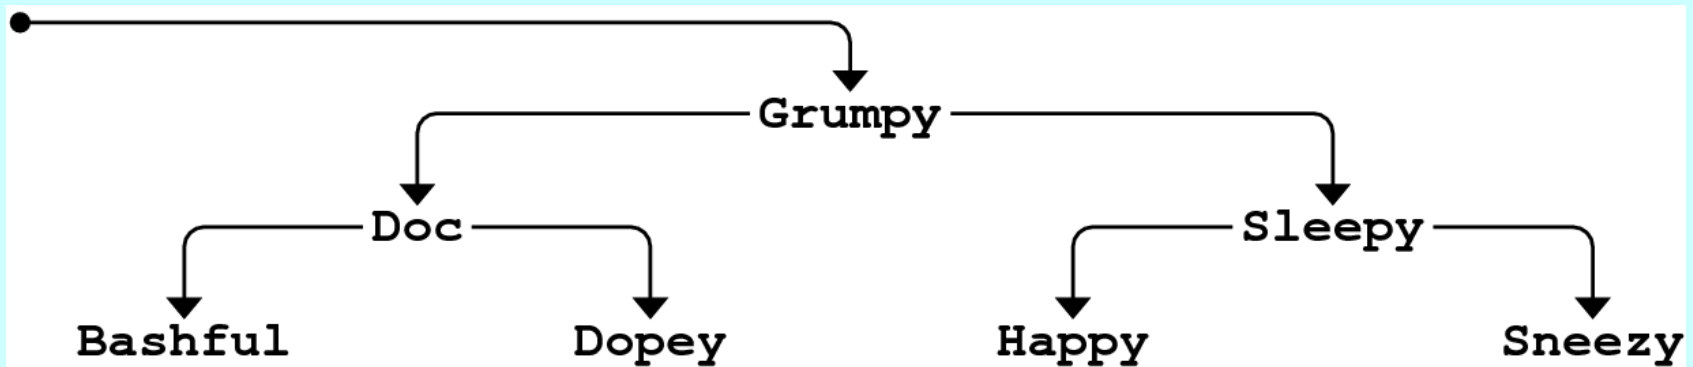
# Exercise: Preorder Traversal




```
void preorderTraversal(Node *t) {  
    if (t != null) {  
        cout << t->key << endl;  
        preorderTraversal(t->left);  
        preorderTraversal(t->right);  
    }  
}
```



# Exercise: Inorder Traversal




```
void inorderTraversal(Node *t) {  
    if (t != null) {  
        inorderTraversal(t->left);  
        cout << t->key << endl;  
        inorderTraversal(t->right);  
    }  
}
```




*An inorder traversal of a BST will always process the nodes in alphabetical order. This is one of the reasons to choose tree-based map over hash map.*





# Exercise: Postorder Traversal



```
void postorderTraversal(Node *t) {  
    if (t != null) {  
        postorderTraversal(t->left);  
        postorderTraversal(t->right);  
        cout << t->key << endl;  
    }  
}
```



# Exercise: Level-order Traversal



*Although the previous traversals process the nodes in different orders, they all reach the nodes in the same order (depth first). Level-order traversal, however, reaches the nodes in a different order (breadth first).*



# Priority Queues

- Besides implementing maps, trees come up in many other programming contexts.
- A **priority queue** is a queue in which the order in which elements are dequeued depends on a numeric priority.
- They are essential to both Dijkstra's and Kruskal's algorithms (examples of some very important **graph** algorithms), therefore, making them efficient improves performance.
- If you implement them in the obvious way using either arrays or linked lists, priority queues require  $O(N)$  time (Exercise 8 from Chapter 14).
- The standard algorithm for implementing priority queues uses a data structure called a **partially ordered tree**, which makes it possible to implement priority queue operations in  $O(\log N)$  time (Exercise 13 from Chapter 16).



# Partially Ordered Trees

- A partially ordered tree is a special class of binary tree (but not a binary search tree) with these additional properties:
  - The tree is **complete**, i.e., it is completely balanced, and each level of the tree is **filled as far to the left as possible**.
  - The root node of the tree has higher priority than the root of either of its subtrees, which are also partially ordered trees.
- A **heap** is an array-based data structure that simulates the operation of a partially ordered tree. (The heap data structure bears no relationship to the pool of unused memory “heap” available for dynamic allocation.)
- The heap stores the nodes in a partially ordered tree simply by counting off the nodes, level by level, from left to right, making it simple to implement tree operations:

<b>parentIndex</b> (n)	is always given by $(n - 1) / 2$
<b>leftChildIndex</b> (n)	is always given by $2 * n + 1$
<b>rightChildIndex</b> (n)	is always given by $2 * n + 2$

# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11


17

The heap:


17					
----	--	--	--	--	--

# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:






# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11

Dequeue the top priority element  $\rightarrow$  11




The heap:




# Exercise: Partially Ordered Trees

Insert in order: 17 93 20 42 68 11

Dequeue the top priority element  $\rightarrow$  11



The heap:



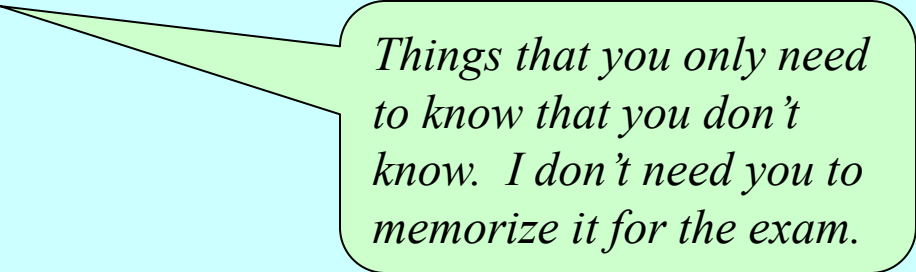
# Heapsort

- Heapsort is a **comparison-based** sorting algorithm utilizing the **heap** data structure.
- Heapsort can be thought of as an **improved selection sort**: like selection sort, heapsort divides its input into a sorted and an unsorted region, and it iteratively shrinks the unsorted region by extracting the largest element from it and inserting it into the sorted region.
- Unlike selection sort, heapsort does not waste time with a linear-time scan of the unsorted region; rather, heapsort maintains the unsorted region in a heap data structure to more quickly find the largest element in each step.
- Although somewhat slower in practice on most machines than a well-implemented quicksort, it has the advantage of a more favorable **worst-case  $O(M\log N)$**  runtime.
- Heapsort is an **in-place** algorithm, but it is not a **stable** sort.



# Heapsort

- Heapify: build an initial heap from the input list.
  - bottom-up SiftDown ( $O(N)$ )
  - top-down SiftUp ( $O(M\log N)$ )
- Swap the first (largest) element with the last element, and no longer consider the last (largest) element.
- SiftDown the current first element to maintain the heap property.
- Repeat the above two steps.



*Things that you only need to know that you don't know. I don't need you to memorize it for the exam.*

# Heapsort

Sorting complexity	Best	Average	Worst
Selection sort	$N^2$	$N^2$	$N^2$
Insertion sort	$N$	$N^2$	$N^2$
Bubble sort	$N$	$N^2$	$N^2$
Merge sort	$N \log N$	$N \log N$	$N \log N$
Heap sort	$N \log N$	$N \log N$	$N \log N$
Quicksort	$N \log N$	$N \log N$	$N^2$

The End